

## Repo layout

This layout assumes a monorepo with shared infra, agents, and per-service apps.

```
text
phase1-ecosystem/
  README.md

config/
  env.global.example.yaml
  cache_key_conventions.yaml
  dragonfly-cluster.yaml

hyperlift/
  creditx-service.yaml
  threat-service.yaml
  guardian-service.yaml
  apps-service.yaml
  phones-service.yaml

agents/
  orchestrator_prompt.yaml
  recovery_agent_prompt.yaml
  agent_registry.yaml

monitoring/
  prometheus-rules.yaml
  grafana-notes.md

services/
  shared/
    python/
      core_cache.py
    node/
      core_cache.ts

  creditx-service/
    app/
      main.py
      db.py
      models.py
      routes_compliance.py
      requirements.txt
      Dockerfile

  threat-service/
    app/
      main.py
      ...
      requirements.txt
      Dockerfile
```

```
guardian-service/
  app/
    main.py
  ...
  requirements.txt
Dockerfile

apps-service/
  src/
    app.ts
    routes.ts
  package.json
  tsconfig.json
Dockerfile

phones-service/
  src/
    app.ts
    routes.ts
  package.json
  tsconfig.json
Dockerfile

.github/
  workflows/
    deploy-creditx.yaml
    deploy-apps.yaml
    deploy-threat.yaml
    deploy-guardian.yaml
    deploy-phones.yaml
```

## Core config & infra files

These files capture the environment contract, Dragonfly cluster, and cache mapping.

`config/env.global.example.yaml`

```
text
# Common ENV (injected via Hyperlift)

# Core infrastructure
SPACES_ENV: "production"
SPACES_REGION: "us-phx-1"

# Database
DATABASE_URL: "postgresql+psycopg2://
ecosystem:***@postgres.internal:5432/ecosystem"

# Cache (Dragonfly)
CACHE_HOST: "dragonfly-cache.internal"
CACHE_PORT: "6379"
CACHE_SSL: "false"
```

```

CACHE_DB_MAIN: "0"          # default; per-service overrides below
CACHE_TIMEOUT_MS: "30000"
CACHE_MAX_POOL_SIZE: "50"

# Observability
PROMETHEUS_ENDPOINT: "http://prometheus.internal:9090"
JAEGER_AGENT_HOST: "jaeger.internal"
JAEGER_AGENT_PORT: "6831"

# Agent mesh
AGENT_REGISTRY_URL: "http://agent-mesh.internal"
ORCHESTRATOR_URL: "http://orchestrator.internal"

# Security
JWT_PUBLIC_KEY: "-----BEGIN PUBLIC KEY-----\n... \n-----END PUBLIC
KEY-----"
OAUTH_ISSUER: "https://auth.ecosystem.ai"
config/cache_key_conventions.yaml

text
cache_key_conventions:
  compliance_documents:
    key: "creditx:doc:{document_id}"
    ttl_sec: 604800           # 7 days

  automation_jobs:
    key: "apps:job:{job_id}"
    ttl_sec: 86400            # 1 day

  threat_events:
    key: "threat:event:{event_id}"
    ttl_sec: 900              # 15 minutes

  device_telemetry:
    key: "guardian:device:{device_id}"
    ttl_sec: 86400            # 1 day

  phone_locations:
    key: "phones:loc:{device_id}"
    ttl_sec: 86400            # 1 day
config/dragonfly-cluster.yaml

text
dragonfly-cluster:
  vm:
    name: cache-prod-01
    tier: Standard-3
    cpu: 4
    ram_gb: 8
    storage_volume:
      name: dragonfly-cache-volume
      size_gb: 100
      mount_path: /mnt/volumes/dragonfly
      encrypted: true          # AES-256

```

```

container:
  image: "dragonflydb/dragonfly:latest"
  ports:
    - 6379
  env:
    DFLY_aof_fsync_sec: 1
    DFLY_max_memory_policy: "allkeys_lru"
    restart_policy: "unless-stopped"
  db_mapping:
    creditx-service:
      db: 0
      prefix: "creditx:"
    threat-service:
      db: 1
      prefix: "threat:"
    guardian-service:
      db: 2
      prefix: "guardian:"
    apps-service:
      db: 3
      prefix: "apps:"
    phones-service:
      db: 4
      prefix: "phones:"
  monitoring:
    prometheus_exporter: true
    alerts:
      - rule: "cache_mem_usage > 0.8"
      - rule: "cache_latency_p95_ms > 10"

```

### **Hyperlift service manifests (Spaceship-native)**

These capture the Spaceship Starlight + Hyperlift deployment for each service from GitHub + Dockerfile.

#### **config/hyperlift/creditx-service.yaml**

```

text
apiVersion: spaceship/v1
kind: Service
metadata:
  name: creditx-service
  namespace: ecosystem-prod
spec:
  source:
    repo: git@github.com:your-org/creditx-service.git
    branch: main
    dockerfile: ./Dockerfile
image:
  repository: spaceship.registry.io/ecosystem/creditx-service

```

```

tag: v2.0.0-dragonfly
replicas:
  min: 3
  max: 10
  target_cpu_utilization: 70
env:
  - name: SPACES_ENV
    value: "production"
  - name: DATABASE_URL
    valueFrom: secretRef:
      name: ecosystem-db
      key: DATABASE_URL
  - name: CACHE_HOST
    value: "dragonfly-cache.internal"
  - name: CACHE_PORT
    value: "6379"
  - name: CACHE_DB
    value: "0"
  - name: CACHE_KEY_PREFIX
    value: "creditx:"
  - name: CACHE_MAX_POOL_SIZE
    value: "50"
health_checks:
  liveness:
    http_get:
      path: /health/live
      port: 8000
  readiness:
    http_get:
      path: /health/ready
      port: 8000
update_strategy:
  type: rolling_update
  max_surge: 1
  max_unavailable: 0
guards:
  error_rate_threshold: 0.01
  latency_p95_ms_threshold: 2000
  auto_rollback: true

```

You can clone this manifest for Threat, Guardian, Apps, and Phones, adjusting CACHE\_DB, CACHE\_KEY\_PREFIX, and repo details.

### Shared cache libraries (Python & Node)

Dragonfly is exposed through Redis-protocol endpoints, so existing Redis clients (`aioredis`, `ioredis`) work with no code-level API change, while benefiting from much higher throughput and lower tail latency.

`services/shared/python/core_cache.py`

```
python
import aioredis
import json
import logging
import os

logger = logging.getLogger(__name__)

class DragonflyCache:
    def __init__(self):
        host = os.getenv("CACHE_HOST", "dragonfly-cache.internal")
        port = os.getenv("CACHE_PORT", "6379")
        db = int(os.getenv("CACHE_DB", "0"))
        self.prefix = os.getenv("CACHE_KEY_PREFIX", "")
        self._dsn = f"redis:///{host}:{port}/{db}"
        self._pool = None

    async def connect(self):
        if self._pool is None:
            self._pool = await aioredis.create_redis_pool(
                self._dsn,
                maxsize=int(os.getenv("CACHE_MAX_POOL_SIZE", "50")),
            )
        logger.info("Connected to Dragonfly cache at %s", self._dsn)

    async def close(self):
        if self._pool:
            self._pool.close()
            await self._pool.wait_closed()

    def _k(self, key: str) -> str:
        return f"{self.prefix}{key}"

    async def get(self, key: str):
        try:
            raw = await self._pool.get(self._k(key))
            if raw is None:
                return None
            return json.loads(raw)
        except Exception as e:
            logger.warning("Cache GET failed for %s: %s", key, e)
            return None

    async def set(self, key: str, value, ttl_sec: int = 3600):
        try:
            raw = json.dumps(value)
```

```

        await self._pool.setex(self._k(key), ttl_sec, raw)
    except Exception as e:
        logger.warning("Cache SET failed for %s: %s", key, e)

async def delete(self, key: str):
    try:
        await self._pool.delete(self._k(key))
    except Exception as e:
        logger.warning("Cache DEL failed for %s: %s", key, e)

async def cache_aside(self, key: str, ttl_sec: int, fetch_fn):
    cached = await self.get(key)
    if cached is not None:
        return cached
    value = await fetch_fn()
    await self.set(key, value, ttl_sec)
    return value
services/shared/node/core_cache.ts

```

```

ts
import Redis from "ioredis";

const host = process.env.CACHE_HOST || "dragonfly-cache.internal";
const port = Number(process.env.CACHE_PORT || 6379);
const db = Number(process.env.CACHE_DB || 3);
const prefix = process.env.CACHE_KEY_PREFIX || "apps:";

const client = new Redis({ host, port, db, lazyConnect: true });

client.on("error", (err) => {
    console.warn("Dragonfly cache error:", err.message);
});

export async function connectCache() {
    if (!client.status || client.status !== "ready") {
        await client.connect();
    }
}

function k(key: string) {
    return `${prefix}${key}`;
}

export async function cacheGet<T>(key: string): Promise<T | null> {
    try {
        const raw = await client.get(k(key));
        return raw ? (JSON.parse(raw) as T) : null;
    } catch (err) {
        console.warn("cacheGet failed:", err);
        return null;
    }
}

export async function cacheSet(key: string, value: any, ttlSec = 3600) {

```

```

    try {
      await client.set(k(key), JSON.stringify(value), "EX", ttlSec);
    } catch (err) {
      console.warn("cacheSet failed:", err);
    }
}

```

## Example service: CreditX (Python/FastAPI)

This is a fully wired example; the other Python services can mirror this pattern, and Node services use the TypeScript cache wrapper.

`services/creditx-service/app/db.py`

```

python
import os
import databases

DATABASE_URL = os.getenv("DATABASE_URL")

database = databases.Database(DATABASE_URL)

```

```

services/creditx-service/app/models.py

python
from pydantic import BaseModel
from typing import Optional

class ComplianceDocument(BaseModel):
  document_id: str
  customer_id: str
  status: str
  payload: dict
  created_at: str
  updated_at: Optional[str]

```

`services/creditx-service/app/routes_compliance.py`

```

python
from fastapi import APIRouter, HTTPException
from .db import database
from .models import ComplianceDocument
from shared.python.core_cache import DragonflyCache
from pybreaker import CircuitBreaker, CircuitBreakerError
import logging

logger = logging.getLogger(__name__)

router = APIRouter(prefix="/creditx", tags=["creditx"])

cache = DragonflyCache()
cache_breaker = CircuitBreaker(fail_max=5, reset_timeout=60)

```

```

@cache_breaker
async def _safe_cache_get(key: str):
    return await cache.get(key)

async def get_with_guard(key: str, fetch_fn, ttl_sec: int):
    try:
        cached = await _safe_cache_get(key)
        if cached is not None:
            return cached
    except CircuitBreakerError:
        logger.warning("Cache circuit OPEN, falling back to DB for %s",
key)
    return await fetch_fn()

    value = await fetch_fn()
    await cache.set(key, value, ttl_sec)
    return value

@router.on_event("startup")
async def startup():
    await database.connect()
    await cache.connect()

@router.on_event("shutdown")
async def shutdown():
    await database.disconnect()
    await cache.close()

@router.get("/documents/{document_id}",
response_model=ComplianceDocument)
async def get_document(document_id: str):
    async def fetch():
        row = await database.fetch_one(
            "SELECT * FROM compliance_documents WHERE document_id"
= :id",
            {"id": document_id},
        )
        if not row:
            raise HTTPException(status_code=404, detail="Document not
found")
        return dict(row)

    key = f"doc:{document_id}"
    result = await get_with_guard(key, fetch, ttl_sec=7 * 24 * 3600)
    return ComplianceDocument(**result)
services/creditx-service/app/main.py

```

```

python
import logging
from fastapi import FastAPI
from .routes_compliance import router as compliance_router

logging.basicConfig(level=logging.INFO)

```

```

app = FastAPI(title="CreditX Service", version="2.0.0-dragonfly")

@app.get("/health/live")
async def live():
    return {"status": "ok"}

@app.get("/health/ready")
async def ready():
    # In production you'd check DB + cache connectivity here
    return {"status": "ready"}
```

**services/creditx-service/requirements.txt**

```

text
fastapi==0.115.0
uvicorn[standard]==0.30.0
databases[postgresql]==0.9.0
aioredis==2.0.1
pybreaker==1.0.1
pydantic==2.9.0
services/creditx-service/Dockerfile
```

```

text
FROM python:3.11-slim

WORKDIR /app

ENV PYTHONUNBUFFERED=1

COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt

COPY app ./app
COPY ../shared/python/core_cache.py ./shared/python/core_cache.py

EXPOSE 8000

CMD ["uvicorn", "app.main:app", "--host", "0.0.0.0", "--port", "8000"]
```

## Apps service (Node.js) – example

**services/apps-service/src/app.ts**

```

ts
import express from "express";
import { connectCache, cacheGet, cacheSet } from "../../shared/node/
core_cache";

const app = express();
app.use(express.json());

app.get("/health/live", (_req, res) => res.json({ status: "ok" }));
```

```

app.get("/health/ready", (_req, res) => res.json({ status: "ready" }));

app.get("/jobs/:id", async (req, res) => {
  const id = req.params.id;
  const key = `job:${id}`;

  await connectCache();

  const cached = await cacheGet<any>(key);
  if (cached) {
    return res.json({ source: "cache", data: cached });
  }

  // TODO: Replace with real DB fetch
  const job = { job_id: id, status: "pending", payload: {} };

  await cacheSet(key, job, 86400);
  return res.json({ source: "db", data: job });
);

export default app;

```

## Agent prompts, registry, and guards

These files encode your Phase 1 agent behavior with Dragonfly-aware semantics.

`config/agents/orchestrator_prompt.yaml`

```

text
system_prompt: |
  You are the Orchestration Agent for the Ecosystem platform.
  The platform is live in production on Spaceship Starlight with
  Dragonfly as the
  shared in-memory cache layer.

  Your job is to coordinate domain agents (CreditX, 91 Apps, Global AI
  Alert,
  Guardian AI, Stolen Phones) to complete workflows reliably.

  Core rules:
  - ALWAYS check dependency readiness (DB, Dragonfly, event bus) before
  dispatching.
  - If Dragonfly is degraded, continue workflows using PostgreSQL and
  mark
    results as "cache_degraded" for observability.
  - Use exponential backoff and circuit breakers to avoid cascading
  failures.
  - Never drop a customer workflow silently; on repeated failure,
  escalate to
    Recovery Agent and queue for manual review.
config/agents/recovery_agent_prompt.yaml

```

```

text
task: |
  When called with a failure event, classify the failure as:
  - transient_network
  - dragonfly_degraded
  - database_issue
  - application_bug
  - configuration_error

  For dragonfly_degraded:
  - Open cache-related circuit breakers.
  - Instruct callers to use database-only code paths.
  - Schedule a health probe against Dragonfly every 30 seconds.
  - When 3 consecutive probes succeed with p95 latency < 5ms, allow
    cache use again.

```

### `config/agents/agent_registry.yaml`

```

text
agent_registry:
  orchestrator-agent:
    url: http://orchestrator.internal
  recovery-agent:
    url: http://recovery.internal
  tuning-agent:
    url: http://tuning.internal

  creditx-compliance-agent:
    url: http://creditx-service.internal
    dependencies: [postgres, dragonfly]

  threat-agent:
    url: http://threat-service.internal
    dependencies: [postgres, dragonfly]

  guardian-agent:
    url: http://guardian-service.internal
    dependencies: [postgres, dragonfly]

  apps-agent:
    url: http://apps-service.internal
    dependencies: [postgres, dragonfly]

  phones-agent:
    url: http://phones-service.internal
    dependencies: [postgres, dragonfly]

```

## CI/CD and deployment guards

Hyperlift still builds from GitHub + Dockerfile and handles blue-green cutovers; here it is wired with pre-flight checks to Dragonfly and readiness endpoints.

```

.github/workflows/deploy-creditx.yaml

text
name: Deploy CreditX (Prod)

on:
push:
  branches: [ main ]
  paths:
    - "services/creditx-service/**"
    - "config/hyperlift/creditx-service.yaml"

jobs:
  build-and-deploy:
    runs-on: ubuntu-latest

    steps:
      - name: Checkout
        uses: actions/checkout@v4

      - name: Set up Docker Buildx
        uses: docker/setup-buildx-action@v3

      - name: Build image
        run: |
          cd services/creditx-service
          docker build -t spaceship.registry.io/ecosystem/creditx-
service:v2.0.0-dragonfly .

      - name: Push image
        run: |
          echo "$REGISTRY_PASSWORD" | docker login spaceship.registry.io
-u "$REGISTRY_USER" --password-stdin
          docker push spaceship.registry.io/ecosystem/creditx-
service:v2.0.0-dragonfly

      - name: Deploy via Hyperlift
        run: |
          spaceship hyperlift deploy \
            --config ../../config/hyperlift/creditx-service.yaml \
            --env production
        env:
          SPACESHIP_TOKEN: ${{ secrets.SPACESHIP_TOKEN }}

config/monitoring/prometheus-rules.yaml

text
groups:
  - name: dragonfly.rules
    rules:
      - alert: DragonflyHighLatency
        expr: dragonfly_request_latency_p95_ms > 10

```

```

for: 5m
labels:
  severity: warning
annotations:
  summary: "Dragonfly p95 latency > 10ms"
  description: "Cache latency elevated, investigate cache or
network."
- alert: DragonflyLowHitRatio
  expr: dragonfly_cache_hit_ratio < 0.5
  for: 10m
  labels:
    severity: warning
  annotations:
    summary: "Dragonfly cache hit ratio < 50%"
    description: "Cache misconfiguration or cold cache, check TTLs
and traffic."
- alert: DragonflyDown
  expr: up{job="dragonfly"} == 0
  for: 1m
  labels:
    severity: critical
  annotations:
    summary: "Dragonfly cache is DOWN"
    description: "Failing over to DB only; investigate
immediately."

```

## How to use this stack

- **For infra:**
  - Provision the Dragonfly VM + container according to `config/dragonfly-cluster.yaml`.
  - Ensure `dragonfly-cache.internal:6379` is reachable from all services.
- **For services:**
  - Wire Python services to `core_cache.py` and Node services to `core_cache.ts`.
  - Ensure env variables in `config/env.global.example.yaml` are provided via Spaceship/Hyperlift secrets.
- **For agents:**
  - Load `orchestrator_prompt.yaml`, `recovery_agent_p`

`rompt.yaml`, and `agent_registry.yaml` into your agent runtime / MCP layer.

- For deployment:
  - Apply Hyperlift manifests in `config/hyperlift/*.yaml`.
  - Enable the GitHub Actions workflows to get auto-deploys on `main`.

With these files in place, Phase 1 is a **Dragonfly-backed, Spaceship-native production build**: Redis is fully removed, Dragonfly is the standard cache layer, and the agents, services, and deployment controls all reflect that reality end-to-end