

DRAFONFLY_72HR_EXECUTION_CHECKLIST.md

DRAFONFLY MIGRATION & DEPLOYMENT CHECKLIST

Jan 16-18, 2026 (72-Hour Execution Plan)

Project: Ecosystem Platform Redis → Dragonfly Migration + Phase 1 Production Go-Live

Timeline: Friday Jan 16 → Sunday Jan 18

Status:  NOT STARTED (ready to go)

JAN 16: DRAGONFLY MIGRATION DAY (Today)

9:00am - 12:00pm: PLANNING & PREPARATION

DevOps Lead - Pre-Flight Checklist:

- Verify Dragonfly Docker image available (`docker pull dragonflydb/dragonfly`)
- Cache VM has 8GB RAM available (check with `free -h`)
- Persistent volume mounted at `/mnt/volumes/dragonfly` (check with `df -h`)
- Network connectivity to cache VM confirmed
- SSH keys configured for all 15 service VMs
- Spaceship CLI installed & configured (`spaceship --version`)
- Docker daemon running on all VMs (`docker ps`)
- PostgreSQL connection verified (`psql -h postgres.internal -U postgres -d ecosystem`)

QA Lead - Test Plan Preparation:

- Staging environment matches production exactly
- Test data loaded into staging PostgreSQL
- Load testing tool installed (`locust` or `jmeter`)
- Monitoring dashboards prepared (Prometheus + Grafana)
- Alert rules configured (latency > 500ms, error rate > 1%)
- Slack channels created:
 - `#ecosystem-dragonfly-migration` (live chat)
 - `#production-incidents` (escalation)
 - `#deployment-status` (updates)

Infrastructure Lead - Risk Mitigation:

- Redis rollback plan documented
- Redis container ready to deploy (standby image)
- Data backup verified (PostgreSQL snapshot created)
- Rollback procedures tested with non-prod environment
- 24-hour rollback window established

Executive Sponsor - Approval:

- Review migration plan (5 min read)
- Approve \$0 investment (migration uses existing infrastructure)
- Authorize go-live decision for Jan 18
- Brief all stakeholders (CTO, CFO, Phase 1 companies)

12:00pm - 3:00pm: DRAGONFLY DEPLOYMENT

DevOps Lead - Cache VM Setup:

1 SSH into cache-prod-01

```

bash
ssh cache-prod-01.ecosystem.internal
2 cd /tmp
3

```

- Connected successfully

4 Stop old Redis container (if running)

```

bash
docker ps | grep redis
5 docker stop redis-old # If exists
6 docker rename redis-old redis-old-$date +%s # Archive it
7

```

- Old Redis stopped & archived

8 Deploy Dragonfly container

```

bash
docker run -d \

```

```
9  --name dragonfly \
10 --restart unless-stopped \
11 --memory 8g \
12 --cpus 4 \
13 -p 6379:6379 \
14 -v /mnt/volumes/dragonfly:/data \
15 -e DFLY_aof_fsync_sec=1 \
16 -e DFLY_max_memory_policy=allkeys_lru \
17 dragonflydb/dragonfly:latest
18
```

- Container started
- Verify: `docker ps | grep dragonfly` (should show RUNNING)

19 Verify Dragonfly connectivity

```
bash
sleep 5 # Wait for startup
20 docker exec dragonfly redis-cli ping
21 # Expected: PONG
22
23 docker exec dragonfly redis-cli INFO server | head -5
24 # Expected: Dragonfly version info
25
26 docker exec dragonfly redis-cli --stat
27 # Expected: Real-time stats
28
```

- Dragonfly responding to Redis commands
- Stats showing 0 keys (fresh instance)

29 Create Dragonfly monitoring script

```
bash
cat > /opt/monitoring/dragonfly-health.sh << 'EOF'
30 #!/bin/bash
31 redis-cli -h cache-prod-01.ecosystem.internal INFO stats
32 redis-cli -h cache-prod-01.ecosystem.internal DBSIZE
33 redis-cli -h cache-prod-01.ecosystem.internal INFO memory
34 EOF
35 chmod +x /opt/monitoring/dragonfly-health.sh
36
```

- Health check script created

Infrastructure Lead - Service Configuration:

1 Update connection strings in all services

```

text
# For each of 5 services, update environment variables:
2 CACHE_URL: redis://cache-prod-01.ecosystem.internal:6379/DB_NUMBER
3 CACHE_TIMEOUT: 30000 # milliseconds
4 CACHE_MAX_POOL_SIZE: 50
5 CACHE_RETRY_COUNT: 3
6

```

- CreditX service: CACHE_URL updated
- Threat service: CACHE_URL updated
- Guardian service: CACHE_URL updated
- Apps service: CACHE_URL updated
- Phones service: CACHE_URL updated

7 Validate connection strings (no typos)

```

bash
# On each service VM:
8 echo $CACHE_URL
9 # Expected: redis://cache-prod-01.ecosystem.internal:6379/X
10

```

- All 5 services have correct URL format

11 Create cache pre-warming script

```

bash
cat > /opt/cache/prewarm.sh << 'EOF'
12 #!/bin/bash
13 redis-cli -h cache-prod-01.ecosystem.internal \
14   --pipe < /opt/cache/initial_keys.txt
15 EOF
16

```

- Pre-warming script ready (for post-deployment)

3:00pm - 6:00pm: INTEGRATION TESTING

QA Lead - Service Integration Tests:

1 Redeploy all 5 services with new connection strings

```

bash
# For each service:
2 spaceship deploy --service creditx --env staging --strategy blue-
green
3 spaceship deploy --service threat --env staging --strategy blue-
green
4 spaceship deploy --service guardian --env staging --strategy blue-
green
5 spaceship deploy --service apps --env staging --strategy blue-
green
6 spaceship deploy --service phones --env staging --strategy blue-
green
7

```

- CreditX deployment complete (health checks passing)
- Threat deployment complete (health checks passing)
- Guardian deployment complete (health checks passing)
- Apps deployment complete (health checks passing)
- Phones deployment complete (health checks passing)

8 Smoke tests for each service

```

bash
# Test 1: Authentication caching
9 curl -X POST https://api-staging.ecosystem.ai/auth/token \
10 -d '{"username": "test", "password": "****"}' \
11 -w "\nStatus: %{http_code}\n"
12 # Expected: 200 (token issued, cached)
13
14 # Test 2: Compliance document caching
15 curl -X POST https://api-staging.ecosystem.ai/creditx/validate \
16 -d '{"document_id": "test-001"}' \
17 -w "\nStatus: %{http_code}\n"
18 # Expected: 200 (cached within 1 second)

```

```
19
20 # Test 3: Threat detection caching
21 curl -X GET https://api-staging.ecosystem.ai/threat/cache/status \
22 -w "\nStatus: %{http_code}\n"
23 # Expected: 200 (cache operational)
24
```

- Auth caching working (1-hour TTL)
- Compliance caching working (7-day TTL)
- Threat caching working (15-min TTL)
- Guardian caching working (24-hour TTL)
- Phones caching working (24-hour TTL)

25 Performance baseline tests

```
bash
# Test 100 concurrent users

26 locust -f /opt/tests/locust_tasks.py \
27 --users 100 \
28 --spawn-rate 10 \
29 --run-time 5m \
30 --host=https://api-staging.ecosystem.ai
31
```

- Response time P95 < 200ms
- Error rate = 0%
- Throughput > 1K req/sec

32 Cache hit ratio measurement

```
bash
# Collect Prometheus metrics

33 curl 'http://prometheus:9090/api/v1/query' \
34 --data-urlencode 'query=rate(cache_hits[5m])'
35
```

- Cache hit ratio recorded (baseline for comparison)

- Expected: ~70% for warm cache

6:00pm - 9:00pm: BLUE-GREEN DEPLOYMENT TO STAGING

DevOps Lead - Staging Deployment:

1 Trigger blue-green deployment

```

bash
# Deploy all services to blue environment (new, with Dragonfly)

2 for service in creditx threat guardian apps phones; do
3   spaceship deploy --service $service \
4     --env staging \
5     --strategy blue-green \
6     --wait-ready 300
7 done
8

```

- All 5 services deployed to blue
- Health checks passing (200 OK)
- Readiness probes green

9 Traffic shift (slow rollout)

```

bash
# Shift 10% traffic to blue

10 spaceship traffic-shift --service all --blue 10 --green 90
11 sleep 300 # Wait 5 minutes
12
13 # Check error metrics
14 curl 'http://prometheus:9090/api/v1/query' \
15   --data-urlencode 'query=rate(errors_total[5m])'
16 # Expected: < 0.1% (near zero)
17
18 # Shift 50% traffic
19 spaceship traffic-shift --service all --blue 50 --green 50
20 sleep 300
21
22 # Shift 100% traffic
23 spaceship traffic-shift --service all --blue 100 --green 0
24

```

- 10% shift successful (5 min stable)

- 50% shift successful (5 min stable)
- 100% shift successful (full traffic on blue)

25 Validation dashboard review

```

bash
# Open Grafana dashboard
26 open https://monitoring-staging.ecosystem.ai:3000
27
28 # Verify metrics:
29 # - Cache hit ratio: > 70%
30 # - Latency P95: < 200ms
31 # - Error rate: < 1%
32 # - Memory usage: < 8GB
33 # - CPU usage: < 60%
34

```

- All metrics within normal range
- Dragonfly running healthy
- No alerts firing

35 Cleanup old environment (optional, can keep 24h)

- Green environment deprecated (kept for rollback)
- Old Redis container archived

End of Day 1 (Jan 16) Status:

```

text
✓ Dragonfly deployed and running
✓ All 5 services connected to Dragonfly
✓ Integration tests passing (100%)
✓ Staging deployment successful
✓ Performance baseline established (70% cache hit ratio)
✓ No errors or warnings
🟡 Ready for break testing (Day 2)

```

JAN 17: BREAK TESTING & CHAOS ENGINEERING (24-HOUR WINDOW)

Failure Scenario Testing

Scenario 1: Dragonfly Container Crash (10am - 11am)

```

bash
# Action: Kill Dragonfly
docker kill dragonfly

```

```

# Expected behavior:
# 1. Services detect connection loss (5 seconds)
# 2. Circuit breaker opens
# 3. Requests rerouted to database (slower, but working)
# 4. Alerts fire in Slack
# 5. On-call engineer paged

# Monitoring:
# - Error rate spikes to 5-10% (expected)
# - Latency increases 10-20x (expected, using DB)
# - Cache hit ratio drops to 0%

# Recovery:
docker start dragonfly
# Expected: Services reconnect within 30 seconds
# Expected: Error rate returns to < 1%
# Expected: Cache hits resume (rebuilding)

# Verification:
# - [ ] Alert fired in Slack (#production-incidents)
# - [ ] Error logs captured
# - [ ] Services recovered automatically
# - [ ] No manual intervention needed

```

Scenario 2: Network Latency (11:30am - 12:30pm)

```

bash
# Action: Add 500ms latency to cache interface
tc qdisc add dev eth0 root netem delay 500ms

# Expected behavior:
# 1. Cache latency increases to ~500ms+
# 2. Circuit breaker activates after 3 timeouts
# 3. Fallback to database (no latency, but slower throughput)
# 4. Alerts: "High latency detected"

# Monitoring:
# - Latency P95: 600-1000ms (high, expected)
# - Error rate: 0-2% (timeout errors)
# - Throughput: Maintained (fallback working)

# Recovery:
tc qdisc del dev eth0 root

```

```

# Verification:
# - [ ] Circuit breaker prevented cascading failures
# - [ ] Fallback mechanism worked
# - [ ] Latency returned to normal within 30s
# - [ ] No data loss

```

Scenario 3: High Concurrency Load (1pm - 2:30pm)

```

bash
# Action: Simulate 500 concurrent users
locust -f locustfile.py --users 500 --spawn-rate 50 --run-time 90m

```

```

# Expected behavior:
# 1. Auto-scaling triggers (more container replicas)
# 2. Connection pool expanded
# 3. Dragonfly memory increases (normal)
# 4. Performance degrades gracefully (P95 < 1000ms acceptable)

# Monitoring:
# - Concurrent connections: 500+ (expected)
# - Memory usage: 4-6GB (normal)
# - CPU: 60-80% (expected under load)
# - Error rate: < 2% (acceptable under stress)
# - Throughput: 10K-15K req/sec

# Verification:
# - [ ] Auto-scaling worked (2+ replicas)
# - [ ] No service crashes
# - [ ] Error rate stayed < 2%
# - [ ] Graceful degradation (not cascading failure)
# - [ ] System recovered after load ended

```

Scenario 4: Partial Cache Failure (3pm - 4pm)

```

bash
# Action: Flush 70% of cache keys
redis-cli -h cache-prod-01.ecosystem.internal FLUSHDB # Clear DB 0
# Manually delete 70% of keys from other DBs

# Expected behavior:
# 1. Cache hit ratio drops to ~30%
# 2. More requests hit database
# 3. Latency increases (database slower than cache)
# 4. System continues working (no errors)
# 5. Tuning agent adjusts cache policy

# Monitoring:
# - Cache hit ratio: 30% (expected, low)
# - Database queries: 3-4x increase (expected)
# - Latency P95: 300-400ms (higher, acceptable)
# - Error rate: 0% (cache misses don't cause errors)

# Recovery:
# Automatic cache repopulation as requests come in

# Verification:
# - [ ] System continued working despite low cache hit ratio
# - [ ] No errors or crashes
# - [ ] Cache rebuilt over time (hit ratio rising)
# - [ ] Database load handled gracefully

```

Scenario 5: Slow Queries (4:30pm - 5:30pm)

```

bash
# Action: Add 100ms delay to all database queries
# (via connection pooling middleware simulation)

```

```
# Expected behavior:  
# 1. Latency increases significantly  
# 2. Circuit breaker evaluates (latency P95 > 200ms)  
# 3. Alerts fire: "High latency - possible DB issue"  
# 4. On-call engineer investigates  
  
# Monitoring:  
# - Latency P95: 500-1000ms (high, expected)  
# - Error rate: 0% (queries still working, just slow)  
# - Throughput: Lower (queuing)
```

```
# Recovery:  
# Remove query delay
```

```
# Verification:  
# - [ ] Alerts fired appropriately  
# - [ ] System remained stable (no crashes)  
# - [ ] Timeout/circuit breaker logic verified  
# - [ ] Recovery automatic
```

Scenario 6: Rollback Test (6pm - 7:30pm)

```
bash  
# Action: Deploy Redis container and switch connection strings  
  
# Step 1: Start Redis  
docker run -d --name redis-old -p 6380:6379 redis:7  
  
# Step 2: Copy data from Dragonfly to Redis (optional)  
redis-cli --rdb /tmp/dragonfly.rdb # Snapshot  
# Or: Use empty cache (let it rebuild)  
  
# Step 3: Update connection strings  
# Change CACHE_URL from :6379 to :6380 (port change)  
  
# Step 4: Restart services  
for service in creditx threat guardian apps phones; do  
  spaceship deploy --service $service --env staging  
done  
  
# Step 5: Verify services working with Redis  
curl https://api-staging.ecosystem.ai/health/ready  
# Expected: 200 OK (all services ready)  
  
# Monitoring:  
# - Services connected to Redis (different port)  
# - API working normally (slower, but functional)  
# - Error rate: 0%  
  
# Step 6: Rollback to Dragonfly  
docker stop redis-old  
# Update connection strings back to :6379  
# Restart services  
for service in creditx threat guardian apps phones; do
```

```
spaceship deploy --service $service --env staging
done
```

```
# Verification:
# - [ ] Seamless switch to Redis (no downtime)
# - [ ] All tests passed with Redis
# - [ ] Seamless switch back to Dragonfly
# - [ ] No data loss during rollback
# - [ ] Performance returned to normal
```

End of Day 2 (Jan 17) Status:

text

```
✓ Scenario 1: Crash recovery - PASSED
✓ Scenario 2: Network latency - PASSED
✓ Scenario 3: High concurrency - PASSED
✓ Scenario 4: Partial cache failure - PASSED
✓ Scenario 5: Slow queries - PASSED
✓ Scenario 6: Rollback test - PASSED
✓ 24-hour stress test complete
✓ Ready for production go-live
```

JAN 18: PRODUCTION GO-LIVE (Phase 1 Companies)

9:00am - 12:00pm: PRODUCTION DEPLOYMENT

DevOps Lead - Production Deployment:

1 Pre-deployment checks

```
bash
# Verify staging is stable (after 24h break testing)
2 curl https://api-staging.ecosystem.ai/health/live # Expected: 200
3 curl https://api-staging.ecosystem.ai/health/ready # Expected:
200
4
5 # Check Dragonfly status on production cache VM
6 ssh cache-prod-01.internal
7 docker ps | grep dragonfly # Expected: RUNNING
8 redis-cli ping # Expected: PONG
9 redis-cli DBSIZE # Expected: (integer) 0 (fresh, will populate)
10
11 # Verify production load balancer ready
12 spaceship lb status --name ecosystem-api-lb # Expected: ACTIVE
13 spaceship lb test --name ecosystem-api-lb # Expected: OK
14
```

- Staging stable (100% uptime, 0% errors)
- Production Dragonfly ready

- Production load balancer ready
- All systems green

15 Blue-green deployment to production

```

bash
# Verify current production is green (stable)
16 spaceship environment status --env production
17 # Expected: all pods running, health check 200 OK
18
19 # Deploy to blue (new environment with Dragonfly)
20 for service in creditx threat guardian apps phones; do
21   spaceship deploy --service $service \
22     --env production \
23     --strategy blue-green \
24     --wait-ready 300
25 done
26

```

- All services deployed to blue
- Health checks passing (200 OK)
- Ready probes green

27 Smoke tests on blue (production)

```

bash
# Test 1: Authentication
28 curl -X POST https://api.ecosystem.ai/auth/token \
29   -H "X-Blue-Green: blue" \
30   -d '{"username": "test", "password": "****"}'
31 # Expected: 200 OK
32
33 # Test 2: Compliance validation
34 curl -X POST https://api.ecosystem.ai/creditx/validate \
35   -H "X-Blue-Green: blue" \
36   -d '{"document_id": "doc-prod-001"}'
37 # Expected: 200 OK
38
39 # Test 3: Threat detection
40 curl -X GET https://api.ecosystem.ai/threat/status \
41   -H "X-Blue-Green: blue"
42 # Expected: 200 OK
43

```

- Auth working on blue
- CreditX working on blue
- Threat detection working on blue
- Guardian working on blue
- Phones working on blue

44 Production traffic shift

```

bash
# Shift 10% traffic to blue (production)

45 spaceship traffic-shift --env production --blue 10 --green 90
46 sleep 300 # Monitor for 5 minutes
47
48 # Check errors
49 curl 'http://prometheus:9090/api/v1/query' \
50   --data-urlencode 'query=rate(errors_total[5m])'
51 # Expected: < 0.1% error rate (near zero)
52
53 # Shift 50%
54 spaceship traffic-shift --env production --blue 50 --green 50
55 sleep 300
56
57 # Shift 100%
58 spaceship traffic-shift --env production --blue 100 --green 0
59

```

- 10% shift: stable (0% errors for 5 min)
- 50% shift: stable (0% errors for 5 min)
- 100% shift: stable (full traffic on blue/Dragonfly)

60 Production metrics validation

```

bash
open https://monitoring.ecosystem.ai:3000

61
62 # Verify metrics:
63 # Dashboard: "Ecosystem Production Dragonfly"
64 # - Cache hit ratio: building toward 70%
65 # - Latency P95: < 200ms
66 # - Error rate: 0%
67 # - Memory: 1-3GB (building)
68 # - CPU: 20-40% (light load, morning)
69 # - Connections: 50-100

```

- All metrics within target range
- No alerts firing
- Dragonfly memory growing (normal, rebuilding cache)

12:00pm - 12:30pm: PHASE 1 COMPANY GO-LIVE NOTIFICATIONS

Customer Success Lead - Stakeholder Communication:

1 Send go-live notification to 5 Phase 1 companies

```

text
Subject: Ecosystem Platform – Phase 1 Go-Live (Jan 18, 2026)

2
3 Dear Leadership,
4
5 The Ecosystem platform is now live in production!
6
7  What's live:
8 • CreditX compliance automation (Revau)
9 • 91 Apps business automation (Nuvei, Master Group)
10 • Global AI Alert threat detection (Spectrum, Comm Tower)
11 • Guardian AI endpoint security (Spectrum)
12 • Stolen Lost Phones device recovery (Spectrum)
13
14  Early metrics (will be available in your dashboards within 1
hour):
15 • Cache performance: 25x faster than industry baseline
16 • API availability: 99.5% SLA
17 • Response times: < 200ms average
18
19  Access your dashboards:
20 • Nuvei: https://dashboard.ecosystem.ai/nuvei
21 • Revau: https://dashboard.ecosystem.ai/revau
22 • Spectrum: https://dashboard.ecosystem.ai/spectrum
23 • Master Group: https://dashboard.ecosystem.ai/master
24 • Comm Tower: https://dashboard.ecosystem.ai/commtower
25
26  Support available 24/7 at support@ecosystem.ai
27
28 Welcome to Phase 1 of the Ecosystem! 
29

```

- Notification sent to Nuvei CEO + CFO

- Notification sent to Revau CEO + CIO
- Notification sent to Spectrum CEO + CIO
- Notification sent to Master Group CEO + Operations
- Notification sent to Comm Tower CEO + Network Ops

30 Enable real-time dashboards for Phase 1 companies

```

bash
# Verify dashboards are live
31 curl https://api.ecosystem.ai/dashboard/nuvei/ebitda
32 # Expected: Real-time EBITDA impact tracking
33
34 curl https://api.ecosystem.ai/dashboard/revau/compliance
35 # Expected: Compliance metrics
36

```

- Nuvei dashboard live (lead scoring metrics)
- Revau dashboard live (compliance metrics)
- Spectrum dashboard live (breach detection metrics)
- Master Group dashboard live (supply chain metrics)
- Comm Tower dashboard live (network threat metrics)

12:30pm - 3:00pm: 24/7 OPERATIONS HANDOFF

On-Call Engineer - Production Monitoring:

1 Take ownership of production monitoring

```

bash
# Set up PagerDuty on-call rotation
2 # Create alert rules:
3 # - Error rate > 1%: Page engineer immediately
4 # - Latency P95 > 500ms: Alert in Slack
5 # - Dragonfly memory > 7GB: Alert in Slack
6 # - Cache hit ratio < 50%: Warning in Slack
7
8 # Verify alerting is working
9 echo "Testing alert: artificial error" >> /var/log/ecosystem/
errors.log
10 # Expected: Alert fired in Slack within 1 minute
11

```

- PagerDuty configured with on-call rotation
- Alert rules created (P1, P2, P3 severity)
- Slack channels integrated (#production-incidents)
- Alert test fired successfully

12 Verify runbooks are accessible

```

bash
# Store runbooks in accessible location
13 cat /opt/runbooks/incident-response.md | wc -l
14 # Expected: Comprehensive guide available
15
16 # Create quick reference for on-call engineer
17 cat /opt/runbooks/quick-reference.txt
18 # Expected: Top 5 incidents + resolution steps
19

```

- Runbooks documented & accessible
- Quick reference card printed/available
- Escalation contacts on display

20 Monitor real-time metrics (Phase 1 companies)

```

bash
# Track EBITDA impact in real-time
21 # Check dashboards every 15 minutes for first 2 hours
22
23 # Expected by end of Day 1:
24 # - Nuvei 91 Apps: +$25M/year pace
25 # - Revau CreditX: +$18M/year pace
26 # - Spectrum Guardian AI: +$8M/year pace
27 # - Master Group 91 Apps: +$15M/year pace
28 # - Comm Tower Global AI Alert: +$6M/year pace
29 # Total Phase 1: $72M/year pace
30

```

- EBITDA dashboards tracking
- Metrics aligned with projections

- Business impact validated

3:00pm+: ONGOING MONITORING & OPTIMIZATION

DevOps Lead - Performance Optimization:

1 Monitor cache hit ratio growth

```
bash
# First hour: 0-20% (cache empty)
2 # Hour 1-2: 20-50% (warming up)
3 # Hour 2-6: 50-70% (optimal)
4 # Hour 6+: 70-85% (mature)
5
```

- Cache hit ratio progression tracked
- No performance regressions
- Dragonfly memory usage reasonable (< 5GB)

6 Tuning agent optimization

```
bash
# Tuning agent continuously optimizes:
7 # - Cache TTL values
8 # - Connection pool sizes
9 # - Auto-scaling thresholds
10 # - Load shedding policies
11
12 # Monitor tuning agent logs
13 tail -f /var/log/ecosystem/tuning-agent.log
14
```

- Tuning agent running (auto-optimization active)
- TTLs adjusted based on hit ratio
- Performance improving over time

15 Recovery agent readiness

```
bash
# Verify recovery agent active
16 curl https://api.ecosystem.ai/agents/recovery/status
```

```
17 # Expected: { "status": "ready", "circuits": "normal" }
18
19 # Verify circuit breakers all CLOSED (healthy)
20 curl https://api.ecosystem.ai/circuits/status
21
```

- Recovery agent ready (auto-healing active)
- All circuit breakers CLOSED (healthy state)
- Agents monitoring for failures

End of Day 3 (Jan 18) Status:

text

- ✓ Production blue-green deployment complete
- ✓ All 5 Phase 1 companies online
- ✓ Real-time dashboards showing EBITDA impact
- ✓ 99.5% uptime SLA maintained
- ✓ Cache hit ratio building (20-50% by end of day)
- ✓ Zero errors or critical incidents
- ✓ 24/7 monitoring operational
- ✓ On-call engineer ready for production support

🎉 PHASE 1 GO-LIVE SUCCESSFUL!

Next: Monitor for 7 days (Jan 18-25), then Phase 2 kickoff

SUCCESS CRITERIA (72-Hour Delivery)

✓ Jan 16 (Dragonfly Migration)

- Dragonfly deployed & running
- All 15 services connected
- Integration tests 100% passing
- Staging deployment successful
- Zero data loss

✓ Jan 17 (Break Testing)

- 6 failure scenarios tested & recovered
- Rollback tested & working
- 24-hour stability validated
- 25x performance improvement verified
- All metrics within target range

✓ Jan 18 (Production Go-Live)

- Production blue-green deployment
- 5 Phase 1 companies online
- EBITDA impact dashboards live
- 99.5% uptime maintained
- Zero production incidents
- 24/7 support operational

SIGN-OFF

- **DevOps Lead:** Dragonfly deployment approved & complete
- **QA Lead:** Break testing complete, zero critical issues
- **Infrastructure Lead:** Rollback procedures verified
- **On-Call Engineer:** Production monitoring ready
- **Project Manager:** Phase 1 go-live approved
- **Executive Sponsor:** Production deployment authorized



READY TO EXECUTE - 72 HOURS TO PHASE 1 GO-LIVE

Next Steps:

- 1 Approve this checklist
- 2 Start Day 1 (Jan 16) at 9am MST
- 3 Follow checklist sequentially
- 4 Report status updates to #ecosystem-dragonfly-migration
- 5 Escalate blockers immediately to DevOps Lead

Contact:

- DevOps Lead: +1-780-555-0100
- On-Call Engineer: PagerDuty (@oncall)
- Project Manager: @project-mgr on Slack

Document Status: READY FOR EXECUTION

Timeline: Jan 16-18, 2026 (72 hours)

Target: Phase 1 production go-live by Jan 18 9am MST

ROI: \$72M annualized EBITDA impact begins Jan 18