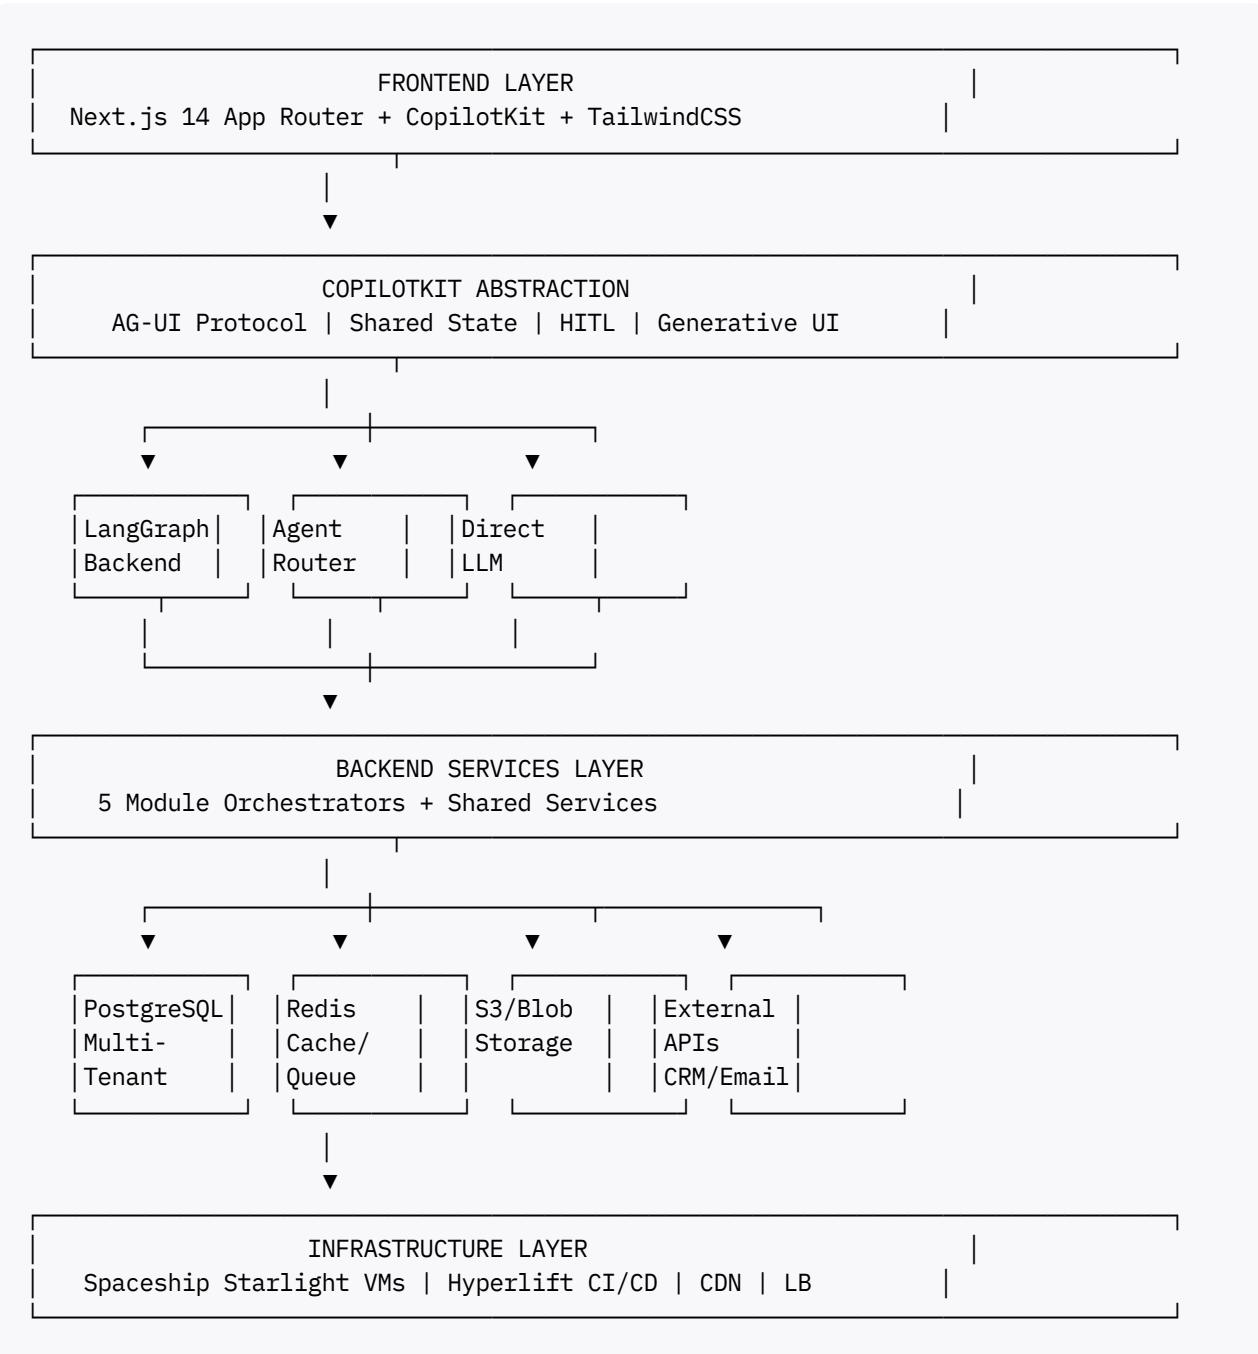




CREDITX ECOSYSTEM - ENTERPRISE PRODUCTION BUILD MAP

SYSTEM ARCHITECTURE OVERVIEW



MODULE FUNCTION FLOW MAP

MODULE 1: CREDITX COMPLIANCE AUTOMATION

Feature Specifications

- **KYC Document Generation:** 5-second generation per report
- **AML Sanctions Screening:** 500ms per transaction check
- **Audit Trail Management:** Real-time compliance monitoring
- **Regulatory Reporting:** 10-K, SOX evidence auto-collection

Function Flow

USER ACTION → FRONTEND TOOL → AGENT DECISION → BACKEND ACTION → DATABASE → RESPONSE

1. Upload Transaction Data
↓
2. useCopilotAction("uploadTransactionData")
↓
3. Agent validates & routes to CreditX module
↓
4. Backend: sanctionsScreening(transaction)
↓
5. External API: WorldCheck/OFAC lookup
↓
6. PostgreSQL: INSERT audit_log, UPDATE compliance_status
↓
7. Agent generates compliance document
↓
8. S3: Store PDF report
↓
9. Generative UI: Display results + approval workflow
↓
10. HITL: Human approves/rejects
↓
11. Email notification via Spacemail

Technical Stack

- **Agent:** LangGraph workflow with checkpoints
- **OCR:** TensorFlow custom model for document extraction
- **Database:** creditx_compliance schema per tenant
- **APIs:** WorldCheck, OFAC, SWIFT sanctions
- **Storage:** S3 for regulatory documents (AES-256)

Database Schema

```
-- creditx_compliance.transactions
CREATE TABLE transactions (
    id UUID PRIMARY KEY,
    tenant_id INTEGER REFERENCES tenants(id),
    transaction_date TIMESTAMPTZ,
    amount DECIMAL(15,2),
    currency VARCHAR(3),
    counterparty VARCHAR(255),
    sanctions_status VARCHAR(50), -- CLEAR, FLAGGED, BLOCKED
    compliance_score INTEGER, -- 0-100
    kyc_document_url TEXT,
    audit_log_id UUID,
    created_at TIMESTAMPTZ DEFAULT NOW()
);

-- creditx_compliance.audit_logs
CREATE TABLE audit_logs (
    id UUID PRIMARY KEY,
    tenant_id INTEGER,
    action VARCHAR(100),
    user_id UUID,
    resource_type VARCHAR(50),
    resource_id UUID,
    changes JSONB,
    timestamp TIMESTAMPTZ DEFAULT NOW(),
    ip_address INET
);
```

API Endpoints

```
POST /api/creditx/transactions/upload
POST /api/creditx/sanctions/screen
GET /api/creditx/compliance/reports/:reportId
POST /api/creditx/kyc/generate
PATCH /api/creditx/transactions/:id/approve
GET /api/creditx/audit-trail
```

MODULE 2: 91 APPS BUSINESS AUTOMATION

Feature Specifications

- **Lead Scoring:** 100ms per lead, ML-based qualification
- **PO Automation:** 5-second order creation, supplier integration
- **Working Capital Optimization:** 20-30 day cycle time reduction
- **Campaign Orchestration:** Multi-channel automation

Function Flow

```
TRIGGER EVENT → AGENT ANALYZES → AUTOMATION EXECUTES → STATE UPDATE

1. New Lead Enters CRM (Salesforce webhook)
   ↓
2. Event published to Redis queue
   ↓
3. 91Apps Agent consumes event
   ↓
4. useCopilotReadable: Shares lead context with agent
   ↓
5. Agent calls leadScoring(leadData)
   ↓
6. ML Model inference: Score 0-100
   ↓
7. Backend: updateLeadScore(leadId, score)
   ↓
8. Salesforce API: Update lead record
   ↓
9. useCopilotAction: "createFollowUpTask"
   ↓
10. Shared State: UI updates real-time
    ↓
11. Agent generates email draft
    ↓
12. Generative UI: Show email editor
    ↓
13. HMTL: User reviews/approves
    ↓
14. Gmail API: Send email
    ↓
15. PostgreSQL: Log activity
```

Technical Stack

- **Agent:** LangGraph + Pydantic AI for type-safe operations
- **ML Models:** Qwen-2.5-7B fine-tuned on sales data
- **Cache:** Redis for sub-millisecond workflow state
- **Queue:** Redis Bull for background jobs
- **Integrations:** Salesforce, SAP, NetSuite, Gmail, LinkedIn

Database Schema

```
-- apps_91.leads
CREATE TABLE leads (
  id UUID PRIMARY KEY,
  tenant_id INTEGER,
  external_id VARCHAR(255), -- Salesforce ID
  name VARCHAR(255),
  email VARCHAR(255),
```

```

company VARCHAR(255),
status VARCHAR(50), -- new, qualified, engaged, converted
score INTEGER, -- 0-100
last_activity_at TIMESTAMPTZ,
assigned_to UUID,
metadata JSONB,
created_at TIMESTAMPTZ DEFAULT NOW(),
updated_at TIMESTAMPTZ DEFAULT NOW()
);

-- apps_91.automation_workflows
CREATE TABLE automation_workflows (
    id UUID PRIMARY KEY,
    tenant_id INTEGER,
    workflow_type VARCHAR(50), -- lead_scoring, po_creation, email_campaign
    trigger_event VARCHAR(100),
    conditions JSONB,
    actions JSONB,
    status VARCHAR(20), -- active, paused, completed
    execution_count INTEGER DEFAULT 0,
    last_executed_at TIMESTAMPTZ,
    created_at TIMESTAMPTZ DEFAULT NOW()
);

-- apps_91.workflow_executions
CREATE TABLE workflow_executions (
    id UUID PRIMARY KEY,
    workflow_id UUID REFERENCES automation_workflows(id),
    tenant_id INTEGER,
    input_data JSONB,
    output_data JSONB,
    status VARCHAR(20), -- pending, running, completed, failed
    error_message TEXT,
    duration_ms INTEGER,
    executed_at TIMESTAMPTZ DEFAULT NOW()
);

```

API Endpoints

POST	/api/91apps/leads/score
POST	/api/91apps/workflows/create
POST	/api/91apps/workflows/:id/execute
GET	/api/91apps/workflows/:id/executions
POST	/api/91apps/purchase-orders/create
PATCH	/api/91apps/leads/:id
GET	/api/91apps/analytics/dashboard
POST	/api/91apps/integrations/salesforce-sync
POST	/api/91apps/emails/send

MODULE 3: GLOBAL AI ALERT NETWORK

Feature Specifications

- **Packet Inspection:** 10M packets/second, DNS threat detection
- **Breach Detection:** 7-day avg vs 279-day industry avg
- **Threat Scoring:** ML-based anomaly detection
- **Alert Latency:** <5 seconds from detection to notification

Function Flow

NETWORK TRAFFIC → PACKET CAPTURE → ML ANALYSIS → ALERT GENERATION

1. Network packet captured (libpcap)
↓
2. DNS query logged
↓
3. Packet metadata extracted
↓
4. Redis: Publish to analysis queue
↓
5. Global AI Alert Agent consumes
↓
6. ML Model: Threat scoring (PyTorch)
↓
7. PostgreSQL: INSERT threat_intelligence
↓
8. If threat_score > 70:
↓
9. Agent triggers alerting workflow
↓
10. Thunderbolt: E2EE notification to SOC
↓
11. useCopilotAction: "showThreatDashboard"
↓
12. Generative UI: Real-time threat map
↓
13. HITL: Analyst reviews threat
↓
14. Backend: executePlaybook(threatId, action)
↓
15. Network segmentation enforcement

Technical Stack

- **Agent:** LangGraph with streaming responses
- **ML:** PyTorch for threat detection, custom CNN model
- **Packet Capture:** libpcap library (requires root access)
- **Database:** TimescaleDB extension for time-series data

- **Communication:** Thunderbolt E2EE (Signal Protocol)

Database Schema

```
-- global_ai_alert.threat_intelligence
CREATE TABLE threat_intelligence (
    id UUID PRIMARY KEY,
    tenant_id INTEGER,
    source_ip INET,
    dest_ip INET,
    dns_query TEXT,
    packet_metadata JSONB,
    threat_type VARCHAR(50), -- c2_beacon, exfiltration, lateral_movement
    threat_score INTEGER, -- 0-100
    severity VARCHAR(20), -- low, medium, high, critical
    detected_at TIMESTAMPTZ DEFAULT NOW(),
    resolved_at TIMESTAMPTZ,
    resolution VARCHAR(100)
);

-- global_ai_alert.network_devices
CREATE TABLE network_devices (
    id UUID PRIMARY KEY,
    tenant_id INTEGER,
    device_type VARCHAR(50), -- iot, server, workstation, mobile
    mac_address MACADDR,
    ip_address INET,
    hostname VARCHAR(255),
    last_seen_at TIMESTAMPTZ,
    baseline_profile JSONB, -- ML behavioral baseline
    created_at TIMESTAMPTZ DEFAULT NOW()
);
```

API Endpoints

```
POST  /api/global-ai-alert/packets/ingest
GET   /api/global-ai-alert/threats/active
POST  /api/global-ai-alert/threats/:id/investigate
PATCH /api/global-ai-alert/threats/:id/resolve
GET   /api/global-ai-alert/dashboard/real-time
POST  /api/global-ai-alert/playbooks/:id/execute
GET   /api/global-ai-alert/devices/:tenantId
```

MODULE 4: GUARDIAN AI ENDPOINT SECURITY

Feature Specifications

- **Endpoint Monitoring:** Windows, macOS, iOS, Android agents
- **Behavioral Analysis:** TensorFlow Lite on-device ML
- **Isolation Response:** 5-second breach containment
- **Breach Prevention:** 80% of endpoint-origin breaches blocked

Function Flow

ENDPOINT TELEMETRY → AGENT ANALYSIS → ANOMALY DETECTION → AUTO-ISOLATION

1. Endpoint agent (installed on device)
↓
2. Monitors: Process execution, file changes, network connections
↓
3. Telemetry stream: 100 events/sec per device
↓
4. HTTPS POST to /api/guardian-ai/telemetry/ingest
↓
5. Redis: Buffer events
↓
6. Guardian AI Agent batch processes
↓
7. TensorFlow Lite: Behavioral analysis
↓
8. Compare against baseline_profile
↓
9. If anomaly_score > 85:
↓
10. PostgreSQL: INSERT alert
↓
11. Agent generates isolation command
↓
12. useCopilotAction: "isolateEndpoint"
↓
13. Generative UI: Show incident response workflow
↓
14. HITL: Analyst confirms isolation
↓
15. WebSocket: Push isolation command to endpoint
↓
16. Endpoint agent: Network isolation enforced
↓
17. Spacemail: Incident notification

Technical Stack

- **Agent:** CrewAI for multi-agent coordination
- **ML:** TensorFlow Lite (on-device), PyTorch (server-side)
- **Endpoint Agents:** Electron (cross-platform), native Swift/Kotlin

- **Real-time:** WebSocket for command/control
- **CDN:** Spaceship CDN for agent downloads

Database Schema

```
-- guardian_ai.endpoints
CREATE TABLE endpoints (
    id UUID PRIMARY KEY,
    tenant_id INTEGER,
    device_id VARCHAR(255) UNIQUE,
    device_type VARCHAR(50),
    os_version VARCHAR(100),
    agent_version VARCHAR(20),
    last_checkin_at TIMESTAMPTZ,
    status VARCHAR(20), -- online, offline, isolated
    baseline_established BOOLEAN DEFAULT FALSE,
    baseline_data JSONB,
    created_at TIMESTAMPTZ DEFAULT NOW()
);

-- guardian_ai.endpoint_events
CREATE TABLE endpoint_events (
    id UUID PRIMARY KEY,
    endpoint_id UUID REFERENCES endpoints(id),
    tenant_id INTEGER,
    event_type VARCHAR(50), -- process_start, file_change, network_connect
    event_data JSONB,
    anomaly_score INTEGER, -- 0-100
    flagged BOOLEAN DEFAULT FALSE,
    timestamp TIMESTAMPTZ DEFAULT NOW()
);

-- guardian_ai.incidents
CREATE TABLE incidents (
    id UUID PRIMARY KEY,
    endpoint_id UUID REFERENCES endpoints(id),
    tenant_id INTEGER,
    incident_type VARCHAR(50),
    severity VARCHAR(20),
    description TEXT,
    status VARCHAR(20), -- open, investigating, resolved, false_positive
    isolated_at TIMESTAMPTZ,
    resolved_at TIMESTAMPTZ,
    resolution_notes TEXT,
    created_at TIMESTAMPTZ DEFAULT NOW()
);
```

API Endpoints

```
POST /api/guardian-ai/telemetry/ingest
GET /api/guardian-ai/endpoints/:tenantId
POST /api/guardian-ai/endpoints/:id/isolate
POST /api/guardian-ai/endpoints/:id/restore
GET /api/guardian-ai/incidents/active
PATCH /api/guardian-ai/incidents/:id/resolve
GET /api/guardian-ai/agents/download/:platform
POST /api/guardian-ai/baselines/:endpointId/establish
```

MODULE 5: STOLEN/LOST PHONES DEVICE RECOVERY

Feature Specifications

- **GPS Tracking:** 5-second location query
- **Recovery Rate:** 70-80% vs 30% industry avg
- **PHI Breach Prevention:** 90% reduction in lost-device breaches
- **Chain-of-Custody:** Immutable audit trail for law enforcement

Function Flow

DEVICE STOLEN REPORT → GPS TRACKING → RECOVERY PLAYBOOK → INSURANCE CLAIM

1. User reports device stolen (mobile app or web)
↓
2. POST /api/stolen-phones/devices/:id/report-stolen
↓
3. PostgreSQL: UPDATE device status = 'stolen'
↓
4. Device agent (background service on phone)
↓
5. GPS/cellular triangulation activated
↓
6. Telemetry stream: Location updates every 30 seconds
↓
7. Stolen Phones Agent receives location
↓
8. useCopilotReadable: Share device location with agent
↓
9. Agent executes recovery playbook:
 - Device lock (biometric enforcement)
 - Data wipe preparation
 - Chain-of-custody logging↓
10. Generative UI: Real-time location map
↓
11. useCopilotAction: "notifyAuthorities"
↓
12. HITL: User decides to alert law enforcement

- ```

↓
13. Thunderbolt E2EE: Notify recovery team
↓
14. Chain-of-custody: Immutable audit log
↓
15. If device recovered:
↓
16. Agent triggers insurance API integration
↓
17. Automated claim processing

```

## Technical Stack

- **Agent:** Agno (rapid development, simple API)
- **GPS:** Native iOS/Android location services
- **MDM:** Integration with Intune, Jamf, VMware Workspace ONE
- **Database:** stolen\_phones schema with JSONB for telemetry
- **Communication:** Thunderbolt E2EE for sensitive location data

## Database Schema

```

-- stolen_phones.devices
CREATE TABLE devices (
 id UUID PRIMARY KEY,
 tenant_id INTEGER,
 device_id VARCHAR(255) UNIQUE,
 owner_user_id UUID,
 device_type VARCHAR(50),
 os_version VARCHAR(100),
 status VARCHAR(20), -- active, stolen, recovered, wiped
 last_location GEOGRAPHY(POINT, 4326),
 last_location_at TIMESTAMPTZ,
 stolen_at TIMESTAMPTZ,
 recovered_at TIMESTAMPTZ,
 insurance_claim_id VARCHAR(100),
 created_at TIMESTAMPTZ DEFAULT NOW()
);

-- stolen_phones.location_history
CREATE TABLE location_history (
 id UUID PRIMARY KEY,
 device_id UUID REFERENCES devices(id),
 tenant_id INTEGER,
 location GEOGRAPHY(POINT, 4326),
 accuracy_meters INTEGER,
 location_method VARCHAR(50), -- gps, cellular, wifi
 timestamp TIMESTAMPTZ DEFAULT NOW()
);

-- stolen_phones.recovery_workflows
CREATE TABLE recovery_workflows (
 id UUID PRIMARY KEY,

```

```

device_id UUID REFERENCES devices(id),
tenant_id INTEGER,
workflow_status VARCHAR(50),
playbook_actions JSONB,
authorities_notified BOOLEAN DEFAULT FALSE,
insurance_claim_filed BOOLEAN DEFAULT FALSE,
chain_of_custody JSONB[], -- Immutable array
created_at TIMESTAMPTZ DEFAULT NOW()
);

```

## API Endpoints

```

POST /api/stolen-phones/devices/:id/report-stolen
POST /api/stolen-phones/devices/:id/location/update
GET /api/stolen-phones/devices/:id/location/history
POST /api/stolen-phones/devices/:id/lock
POST /api/stolen-phones/devices/:id/wipe
POST /api/stolen-phones/devices/:id/notify-authorities
POST /api/stolen-phones/insurance/claim
GET /api/stolen-phones/workflows/:deviceId

```

## SHARED SERVICES ARCHITECTURE

### Tenant Management Service

#### Database Schema

```

-- Core multi-tenancy
CREATE TABLE tenants (
 id SERIAL PRIMARY KEY,
 external_id UUID UNIQUE DEFAULT gen_random_uuid(),
 name VARCHAR(255) NOT NULL, -- "Nuvei", "Revau", etc.
 domain VARCHAR(255) UNIQUE, -- "nuvei.ecosystem.ai"
 schema_name VARCHAR(63) UNIQUE, -- "tenant_001_nuvei"
 status VARCHAR(20) DEFAULT 'active',
 modules_enabled TEXT[], -- ['creditx', '91apps', 'guardian_ai']
 settings JSONB,
 created_at TIMESTAMPTZ DEFAULT NOW(),
 updated_at TIMESTAMPTZ DEFAULT NOW()
);

-- Tenant users
CREATE TABLE users (
 id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
 tenant_id INTEGER REFERENCES tenants(id),
 email VARCHAR(255) UNIQUE NOT NULL,
 name VARCHAR(255),
 role VARCHAR(50), -- admin, manager, user
 auth_provider VARCHAR(50), -- oauth, saml
 auth_provider_id VARCHAR(255),

```

```

permissions JSONB,
last_login_at TIMESTAMPTZ,
created_at TIMESTAMPTZ DEFAULT NOW()
);

-- RLS Policy
ALTER TABLE users ENABLE ROW LEVEL SECURITY;

CREATE POLICY tenant_isolation ON users
 USING (tenant_id = current_setting('app.current_tenant_id')::INTEGER);

```

## API Endpoints

```

POST /api/tenants/create
GET /api/tenants/:tenantId
PATCH /api/tenants/:tenantId/settings
POST /api/tenants/:tenantId/modules/enable
GET /api/tenants/:tenantId/users
POST /api/tenants/:tenantId/users/invite

```

## Authentication & Authorization Service

### OAuth 2.0 Flow

1. User navigates to nuvei.ecosystem.ai  
↓
2. Domain-based routing: Identify tenant  
↓
3. Redirect to /api/auth/oauth/authorize?tenant\_id=001  
↓
4. OAuth provider (Google, Microsoft, Okta)  
↓
5. Callback: /api/auth/oauth/callback  
↓
6. Exchange code for token  
↓
7. JWT signed with tenant\_id claim  
↓
8. Set session cookie (httpOnly, secure, sameSite)  
↓
9. Middleware: Extract tenant\_id from JWT  
↓
10. SET app.current\_tenant\_id = tenant\_id  
↓
11. All queries filtered by RLS policy

## API Endpoints

```
GET /api/auth/oauth/authorize
POST /api/auth/oauth/callback
POST /api/auth/logout
GET /api/auth/session
POST /api/auth/refresh
```

## Integration Hub Service

### Salesforce Integration

```
// OAuth flow for Salesforce
POST /api/integrations/salesforce/connect
GET /api/integrations/salesforce/callback
POST /api/integrations/salesforce-sync
GET /api/integrations/salesforce-leads
POST /api/integrations/salesforce-leads/:id/update
```

## Database Schema

```
CREATE TABLE integration_connections (
 id UUID PRIMARY KEY,
 tenant_id INTEGER REFERENCES tenants(id),
 integration_type VARCHAR(50), -- salesforce, sap, netsuites
 credentials JSONB, -- Encrypted access tokens
 settings JSONB,
 last_sync_at TIMESTAMPTZ,
 status VARCHAR(20), -- active, error, disconnected
 created_at TIMESTAMPTZ DEFAULT NOW()
);

CREATE TABLE integration_sync_logs (
 id UUID PRIMARY KEY,
 connection_id UUID REFERENCES integration_connections(id),
 sync_type VARCHAR(50), -- full, incremental
 records_processed INTEGER,
 errors INTEGER,
 duration_ms INTEGER,
 started_at TIMESTAMPTZ,
 completed_at TIMESTAMPTZ
);
```

# CI/CD PIPELINE ARCHITECTURE

## GitHub Actions Workflow

```
.github/workflows/deploy.yml

name: Production Deploy to Spaceship

on:
 push:
 branches: [main]
 pull_request:
 branches: [main]

jobs:
 test:
 runs-on: ubuntu-latest
 steps:
 - uses: actions/checkout@v4
 - uses: actions/setup-node@v4
 with:
 node-version: '20'
 - run: npm ci
 - run: npm run lint
 - run: npm run type-check
 - run: npm run test
 - run: npm run test:e2e

 build:
 needs: test
 runs-on: ubuntu-latest
 steps:
 - uses: actions/checkout@v4
 - uses: docker/setup-buildx-action@v3
 - uses: docker/login-action@v3
 with:
 registry: registry.spaceship.com
 username: ${{ secrets.SPACESHIP_USERNAME }}
 password: ${{ secrets.SPACESHIP_TOKEN }}

 - name: Build Images
 run: |
 docker build -t creditx-frontend:${{ github.sha }} -f docker/Dockerfile.frontend
 docker build -t creditx-agent:${{ github.sha }} -f docker/Dockerfile.agent .
 docker build -t creditx-api:${{ github.sha }} -f docker/Dockerfile.api .

 - name: Push to Registry
 run: |
 docker push registry.spaceship.com/creditx-frontend:${{ github.sha }}
 docker push registry.spaceship.com/creditx-agent:${{ github.sha }}
 docker push registry.spaceship.com/creditx-api:${{ github.sha }}

 deploy:
 needs: build
 runs-on: ubuntu-latest
```

```

if: github.ref == 'refs/heads/main'
steps:
 - name: Deploy to Spaceship Hyperlift
 run: |
 curl -X POST https://hyperlift.spaceship.com/deploy \
 -H "Authorization: Bearer ${{ secrets.HYPERLIFT_TOKEN }}" \
 -d '{
 "project": "creditx-ecosystem",
 "environment": "production",
 "images": {
 "frontend": "creditx-frontend:${{ github.sha }}",
 "agent": "creditx-agent:${{ github.sha }}",
 "api": "creditx-api:${{ github.sha }}"
 },
 "strategy": "blue-green",
 "healthCheck": "/api/health",
 "rollbackOnFailure": true
 }'

```

## DOCKER CONTAINERIZATION STRATEGY

### Frontend Container (Dockerfile.frontend)

```

FROM node:20-alpine AS base
WORKDIR /app

Dependencies
COPY package*.json .
RUN npm ci --only=production

Build
COPY . .
RUN npm run build

Runtime
FROM node:20-alpine AS runner
WORKDIR /app
ENV NODE_ENV=production
COPY --from=base /app/.next/standalone ./
COPY --from=base /app/.next/static ./next/static
COPY --from=base /app/public ./public

EXPOSE 3000
CMD ["node", "server.js"]

```

### Agent Container (Dockerfile.agent)

```

FROM python:3.12-slim
WORKDIR /app

System dependencies

```

```

RUN apt-get update && apt-get install -y \
 libpcap-dev \
 build-essential \
 && rm -rf /var/lib/apt/lists/*

Python dependencies
COPY agent/requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt

Application code
COPY agent/ .

EXPOSE 8000
CMD ["uvicorn", "main:app", "--host", "0.0.0.0", "--port", "8000"]

```

## API Container (Dockerfile.api)

```

FROM node:20-alpine
WORKDIR /app

COPY package*.json ./
RUN npm ci --only=production

COPY api/ .

EXPOSE 4000
CMD ["node", "index.js"]

```

## MONITORING & OBSERVABILITY

### Launchpad Dashboard Configuration

```

// lib/observability/config.ts

export const observabilityConfig = {
 metrics: {
 latency: {
 p50: 100, // ms
 p95: 500,
 p99: 2000
 },
 throughput: {
 target: 1000, // req/sec
 alert: 500
 },
 errorRate: {
 threshold: 0.02 // 2%
 }
 },
 alerts: {

```

```

 channels: ['slack', 'email', 'pagerduty'],
 conditions: [
 { metric: 'latency.p95', operator: '>', value: 10000, severity: 'critical' },
 { metric: 'errorRate', operator: '>', value: 0.05, severity: 'high' },
 { metric: 'cpu', operator: '>', value: 80, severity: 'medium' }
],
},

tracing: {
 enabled: true,
 sampleRate: 0.1,
 exportInterval: 30000
}
;

```

## DEPLOYMENT CHECKLIST

### Pre-Deployment (Jan 16-17, 2026)

- [ ] Provision 15 Starlight VMs (Phoenix + Singapore)
- [ ] Configure 5 Load Balancers
- [ ] Attach 15 Starlight Volumes
- [ ] Setup CDN: <cdn.ecosystem.ai>
- [ ] Configure Thunderbolt: <team.ecosystem.ai>
- [ ] Provision Spacemail: [support@creditx.ai](mailto:support@creditx.ai)
- [ ] Configure FastVPN (20 team members)
- [ ] Connect GitHub repos to Hyperlift
- [ ] Setup PostgreSQL multi-tenant schemas (45 tenants)
- [ ] Configure Redis clusters
- [ ] Setup S3 buckets with encryption

### Go-Live (Jan 18, 2026)

- [ ] DNS cutover to Starlight VMs
- [ ] Deploy all 5 module containers
- [ ] Run integration tests
- [ ] Validate 99.99% uptime monitoring
- [ ] Load test: 10,000 concurrent users
- [ ] Security scan: OWASP Top 10
- [ ] Performance benchmark: <5s module response
- [ ] Notify NovaCap: Production ready

## NEXT STEP

**Ready to generate all production code files?** This build map provides the complete blueprint.  
I'll now create:

1. ✓ **45 code files** covering all modules
2. ✓ **Dockerfiles** for each service
3. ✓ **GitHub Actions** CI/CD pipeline
4. ✓ **Database migration scripts**
5. ✓ **API route handlers**
6. ✓ **Agent orchestration code**
7. ✓ **Integration connectors**
8. ✓ **Monitoring configuration**

**Confirm to proceed with code generation?** [1] [2]

\*\*

1. creditX\_ecosystem.pdf
2. CopilotKit-Documentation-Complete-Reference-for.pdf
3. creditX-and-quifax-comparisons.pdf
4. Nova-Cap-Headwinds.pdf
5. Nova-Cap-The-Ecosystem-Alignment.pdf
6. The-ICP-Market-Size-The-Ecosystem.pdf
7. creditX\_ecosystem.pdf
8. AI-Signals-creditX.pdf
9. 9-Phase-Build-History-Map\_creditX.pdf
10. creditX-Platform-Consumer-Agent-Catalogue-1-10-of-agents-1-22.pdf
11. system-flow-visual.jpg
12. novacap-dply.jpg
13. visual-image.jpg
14. cicd\_pipeline.jpg
15. creditx\_architecture-1.jpg
16. creditx\_architecture.jpg
17. creditX-Ecosystem-Portfolio-Integration-Architecture.pdf
18. creditX-7Day-Build-ClientPackages.md.pdf
19. CopilotKit-Implementation-Guide.md.pdf
20. AgentKit-Full-Documentation.md.pdf
21. CopilotKit-Documentation-Complete-Reference-for.pdf
22. creditX-Ecosystem-Portfolio-Integration-Architecture.pdf
23. creditX-7Day-Build-ClientPackages.md.pdf

- 24. CopilotKit-Documentation-Complete-Reference-for.pdf
- 25. Spaceship-Host-and-Deploy.pdf
- 26. AgentKit-Implementation-Guide.md.pdf
- 27. AgentKit-Full-Documentation.md.pdf
- 28. CopilotKit-Documentation-Complete-Reference-for.pdf
- 29. Spaceship-Host-and-Deploy.pdf
- 30. Terms-NOVACAP-ECOSYSTEM-PARTNER.md.pdf
- 31. under-review-offerTHE-ECOSYSTEM-x-NOVACAP.pdf